


Federation of Kubernetes Clusters




About Me



Stay in Touch!



@Joshuaaweiwei



@CocoonJoshua

Principal Architect @TenxCloud

<https://github.com/JoshuaAndrew>

<https://blog.csdn.net/chweiweich>

Agenda

1. Federation community current situation
 - 1.1> federation v1
 - 1.2> federation v2
2. Federation problems to be solved urgently
 - 2.1> Global Load balancing
 - 2.2> Global Distributed Storage

Federation

What use cases require Cluster Federation?

1. Capacity Overflow

- location affinity
- cross-cluster scheduling
- cross-cluster service discovery
- cross-cluster migration
- cross-cluster load-balancing
- cross-cluster Data synchronization
- cross-cluster monitoring and auditing

2. Sensitive Workloads

What kinds of rules determine which workloads go where?

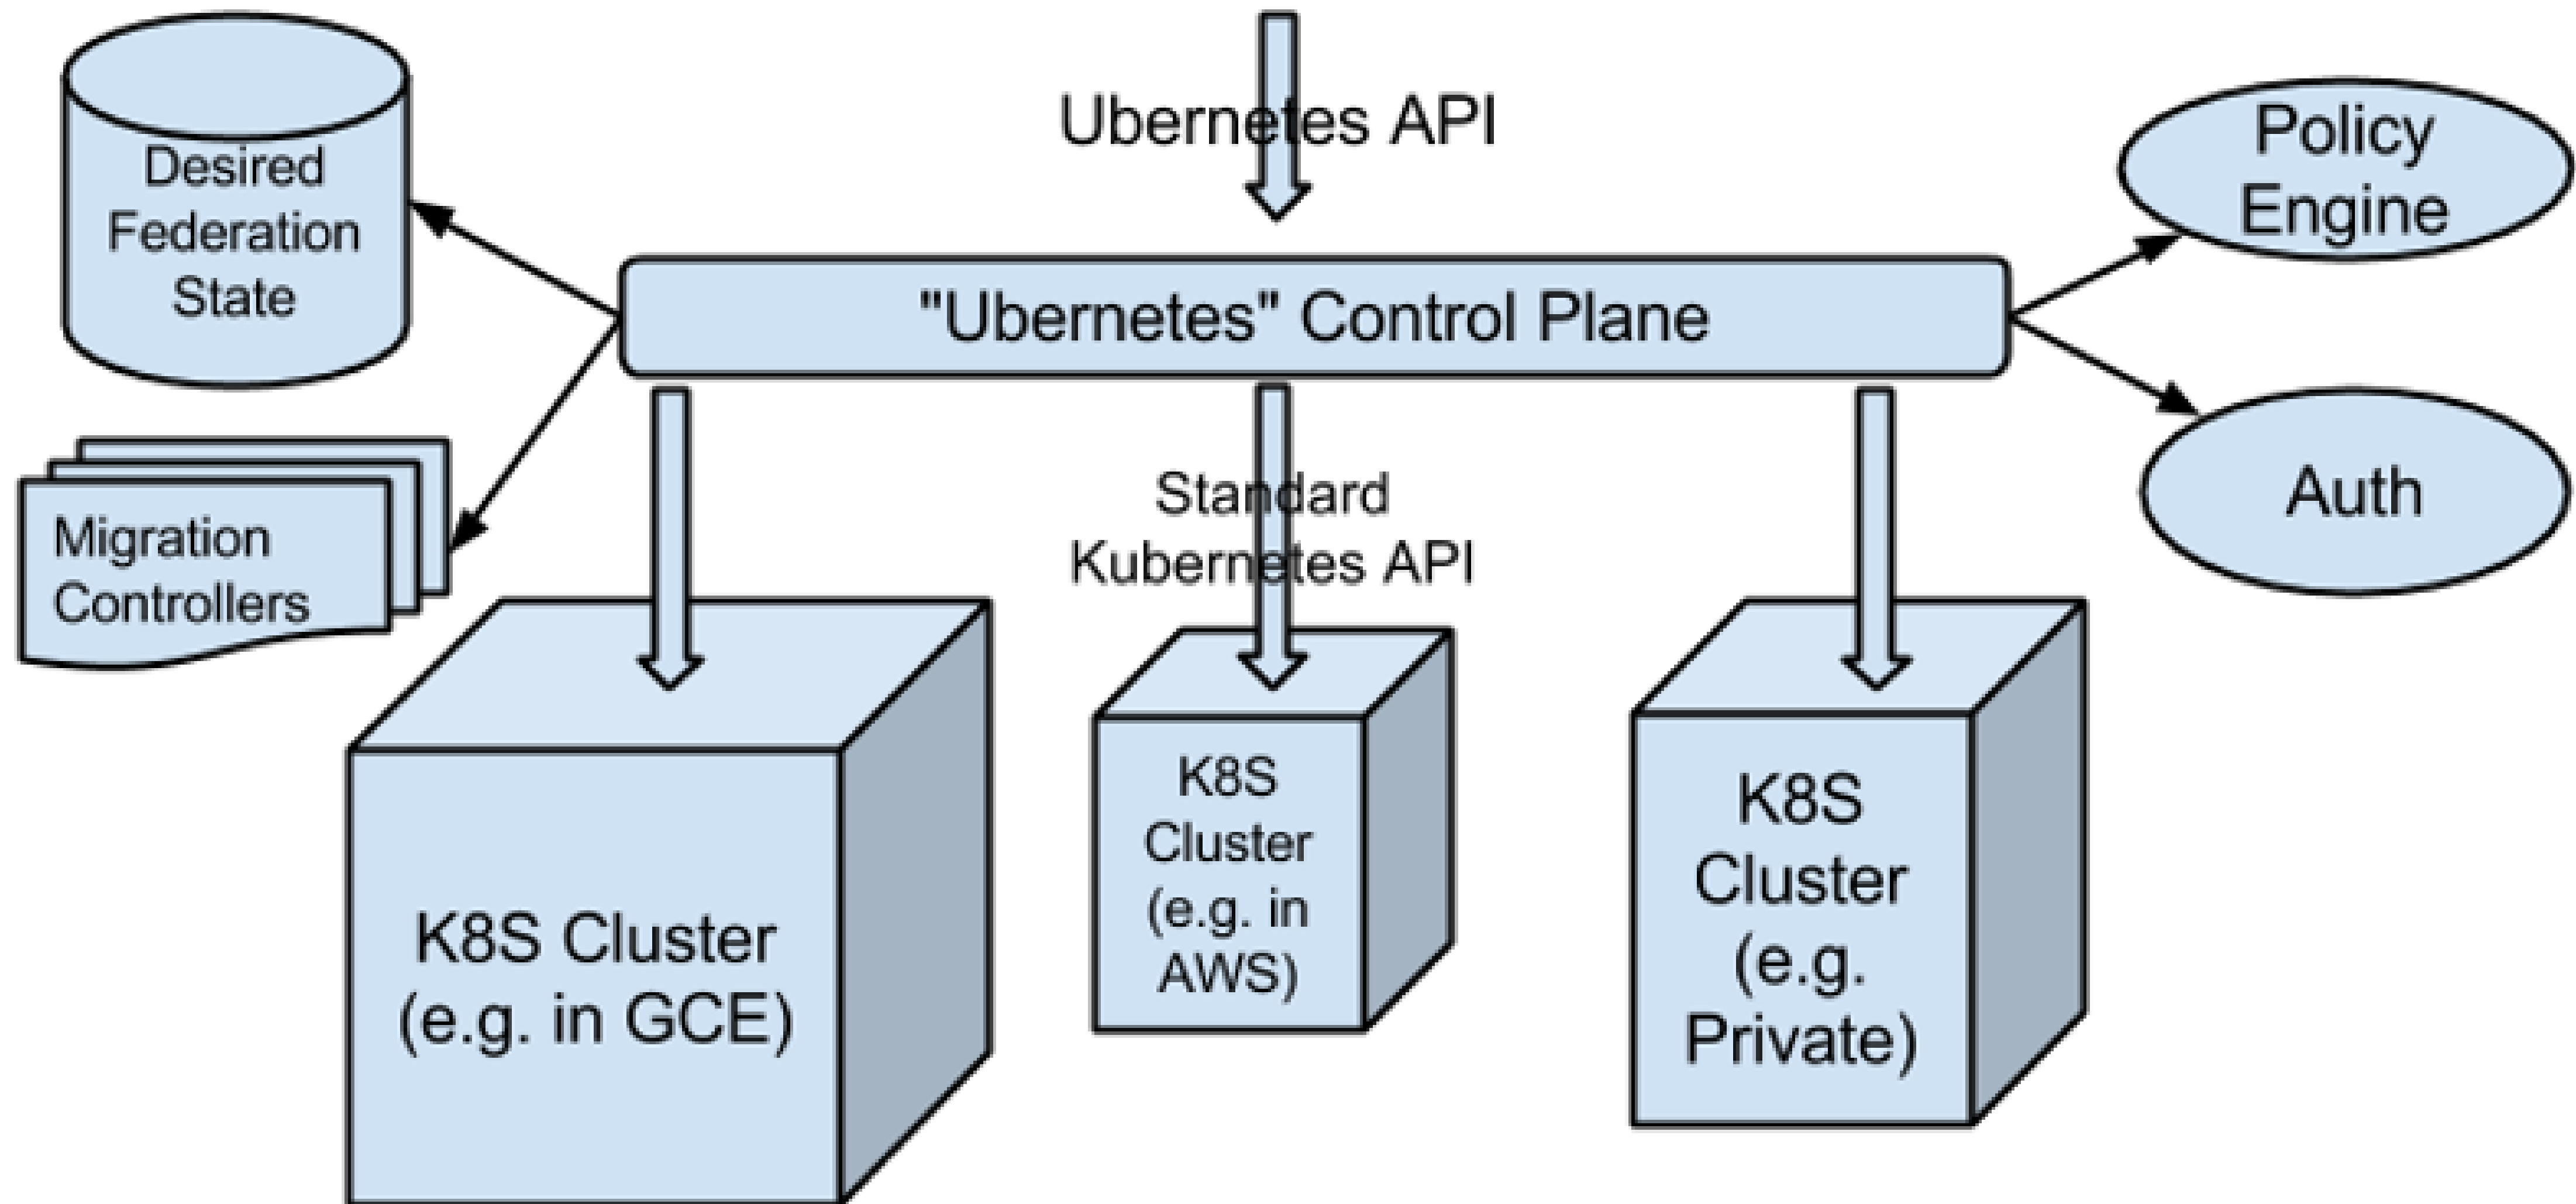
3. Vendor lock-in avoidance

workloads to run across multiple cloud providers at all times

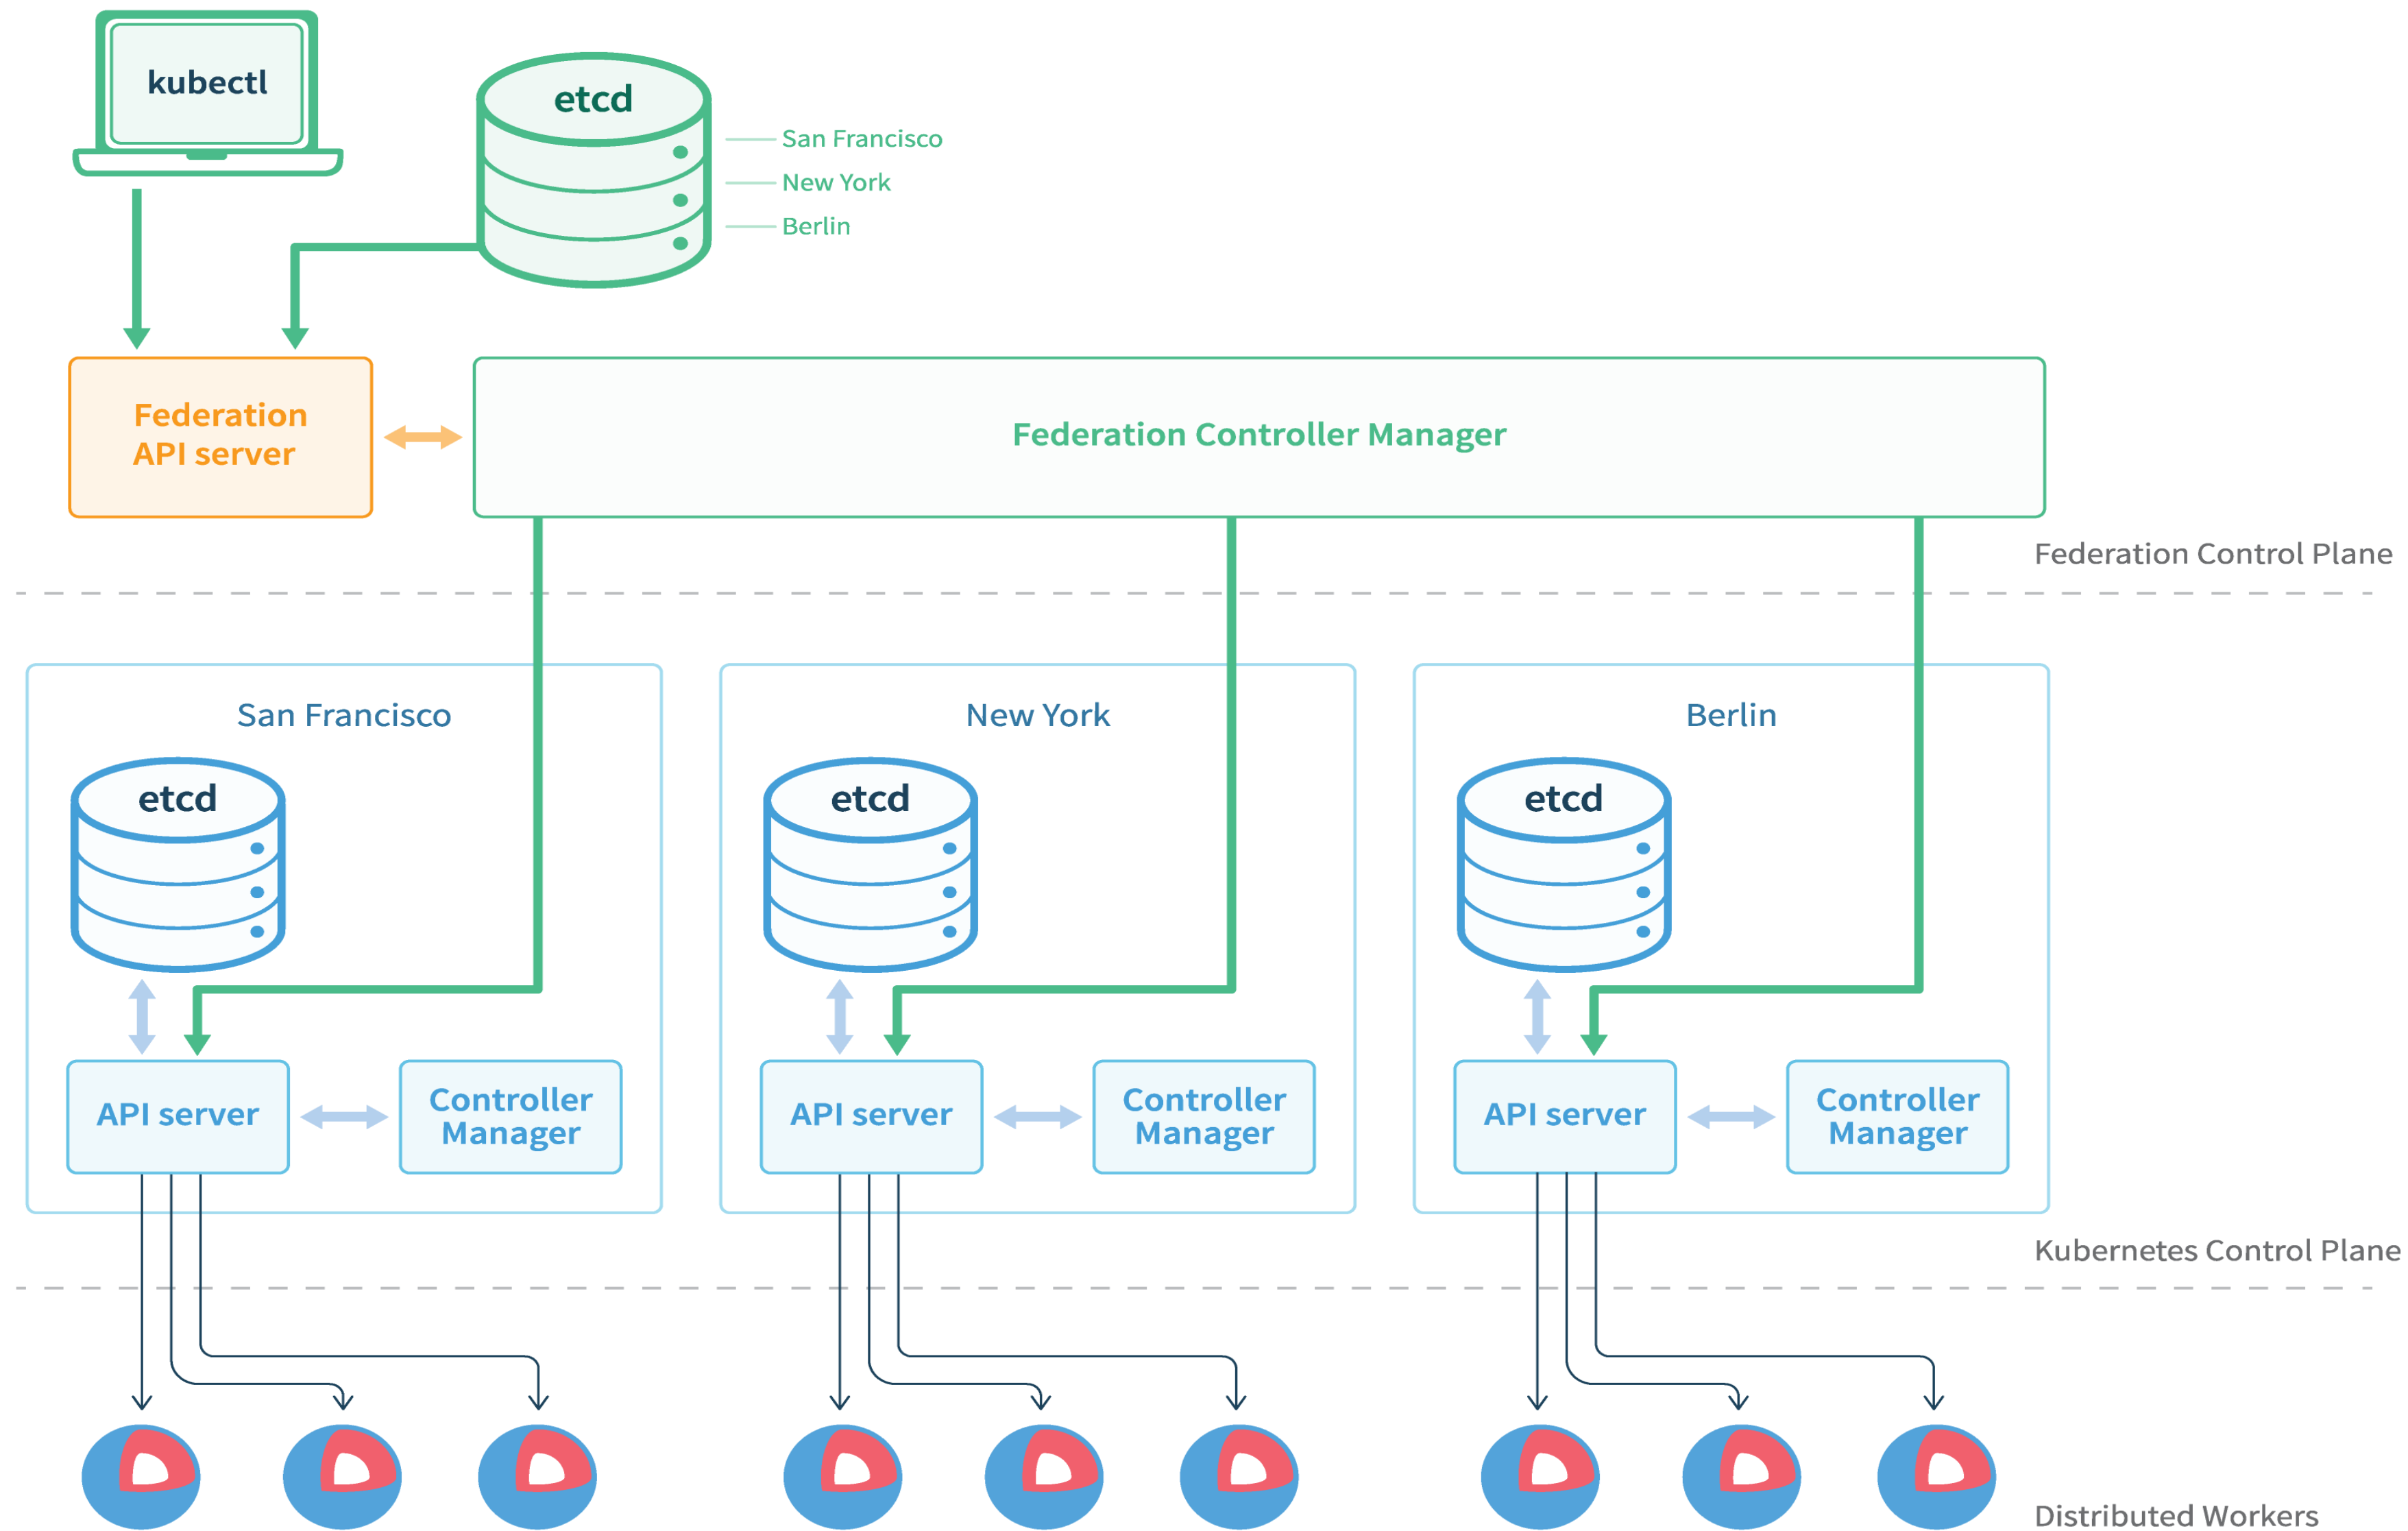
4. High Availability

spread my service across multiple such zones and have my service remain available even if one of the availability zones or cloud providers "goes down"

Federation v1



Federation v1

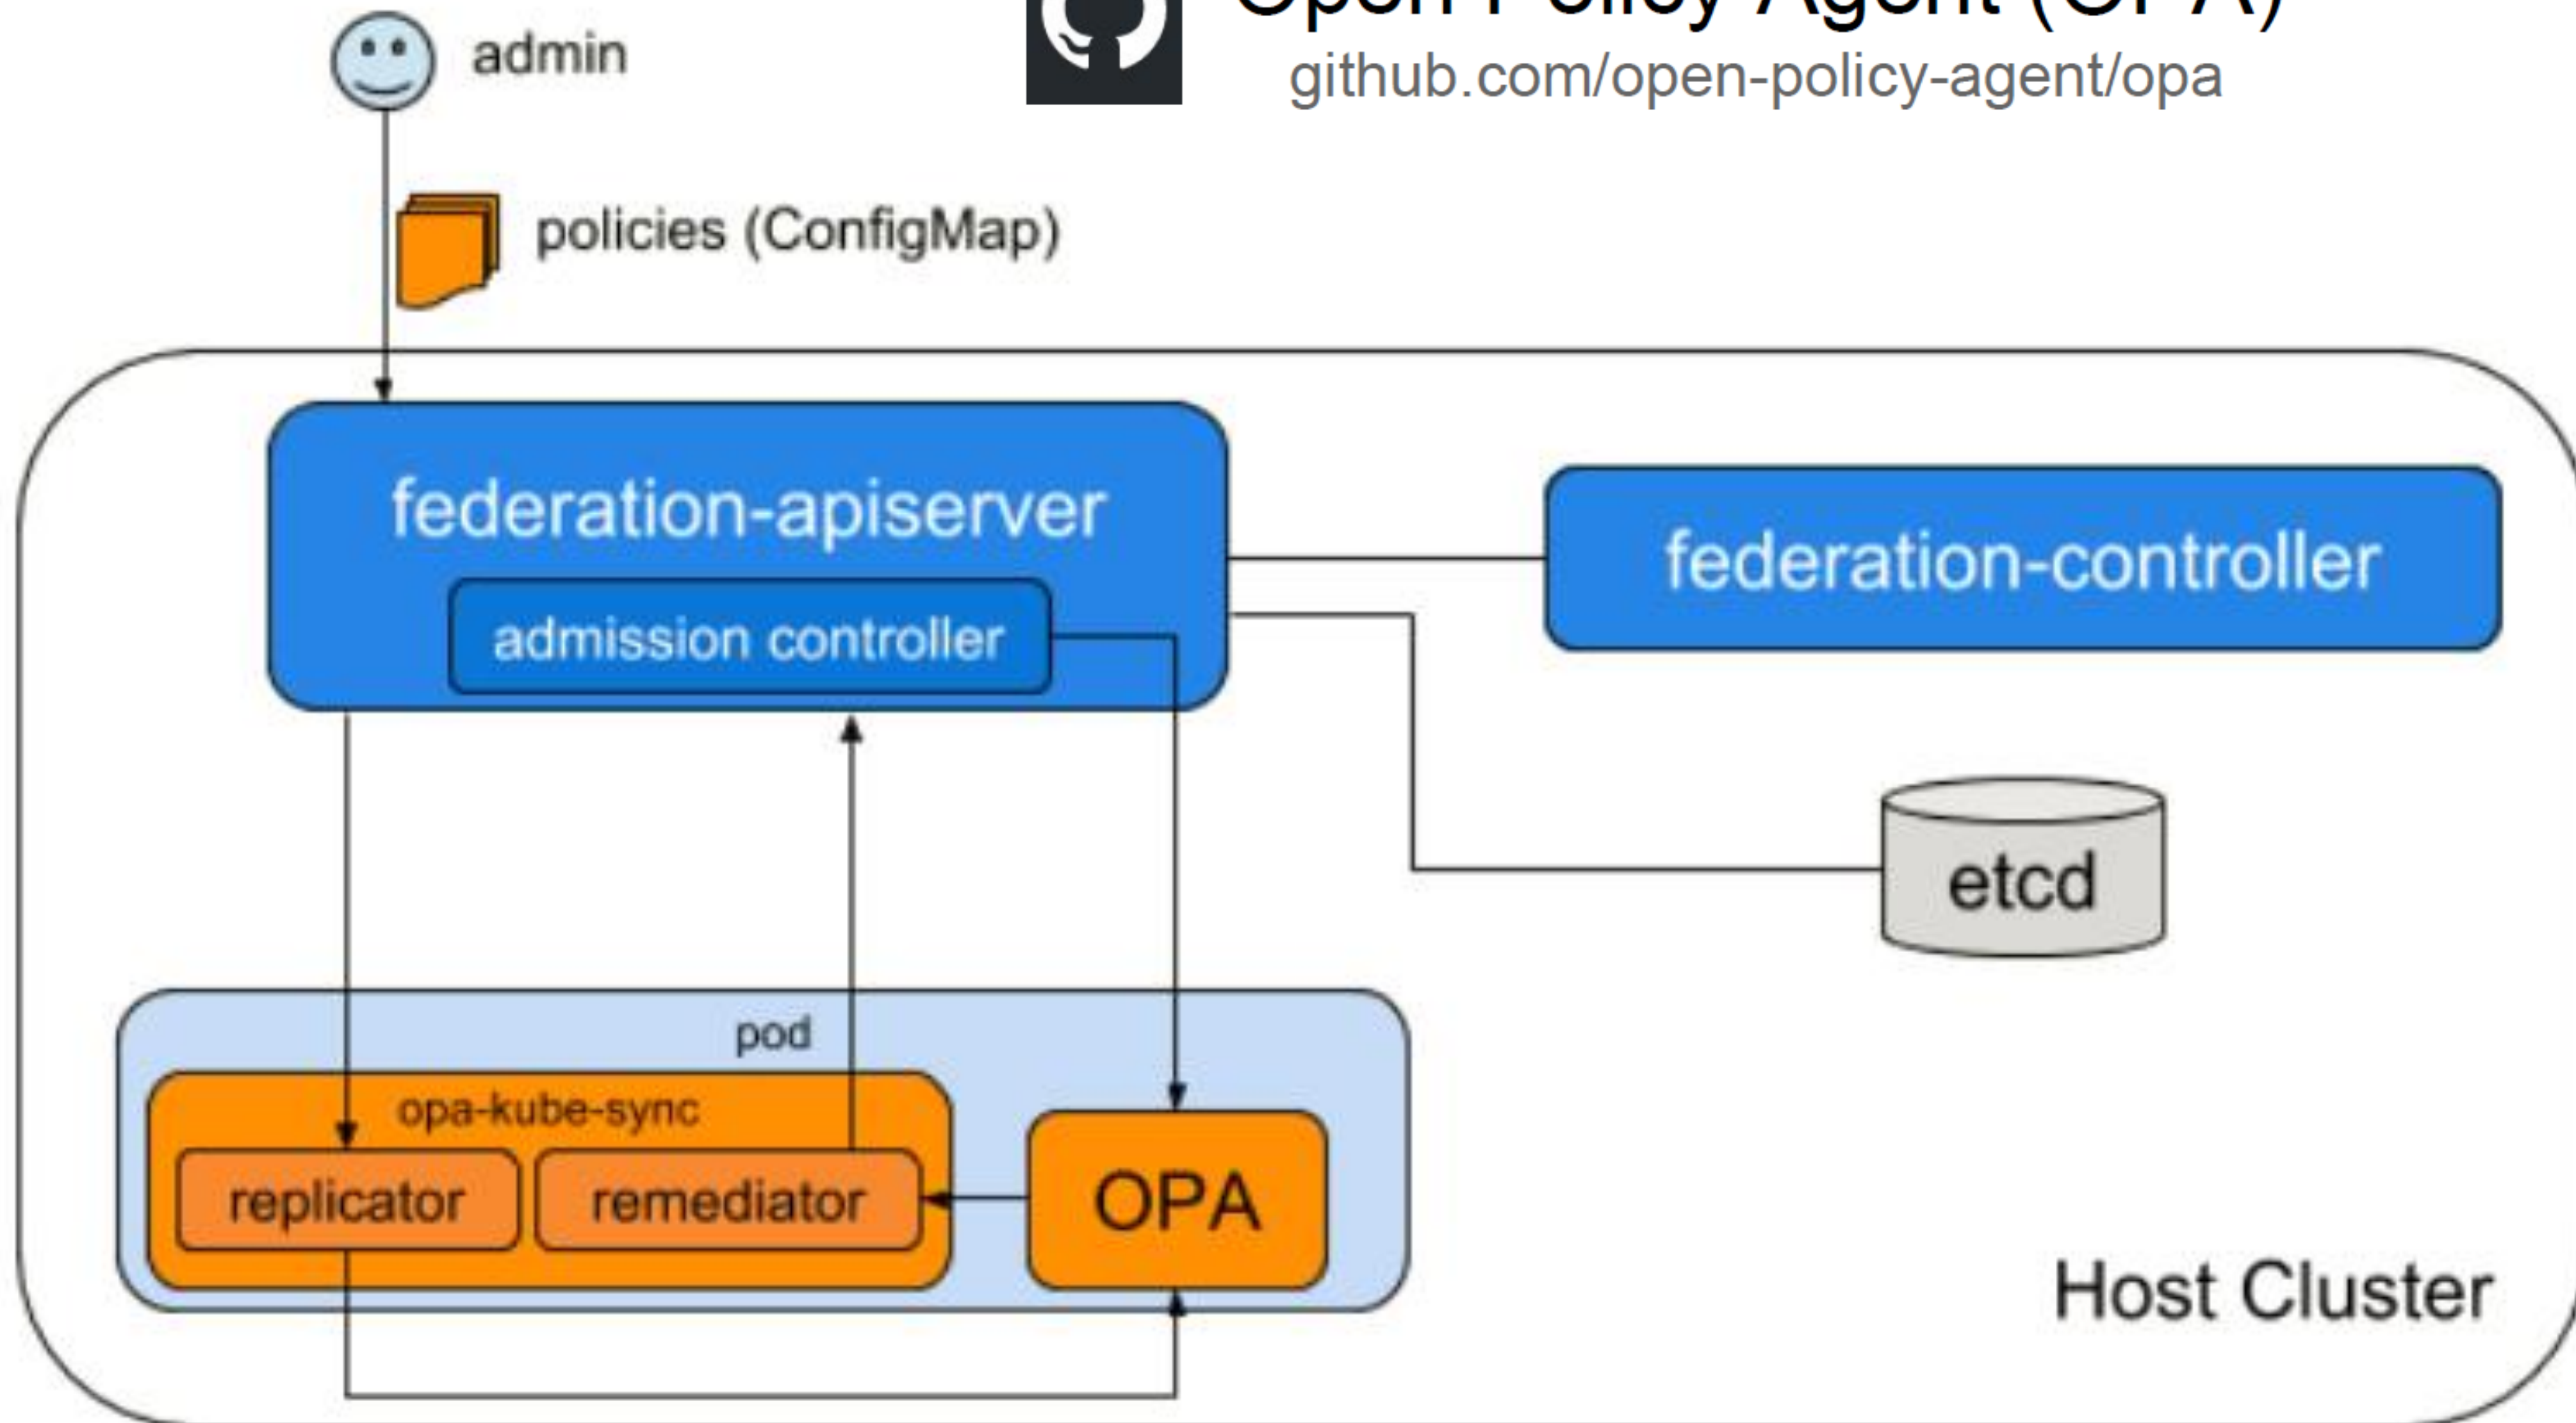


Federation v1



Open Policy Agent (OPA)

github.com/open-policy-agent/opa



Federation v1 -- Resources

Top Level API Objects

- [v1beta1.Deployment](#)
- [v1beta1.DeploymentList](#)
- [v1beta1.DeploymentRollback](#)
- [v1beta1.Scale](#)
- [v1beta1.DaemonSetList](#)
- [v1beta1.DaemonSet](#)
- [v1beta1.Ingress](#)
- [v1beta1.IngressList](#)
- [v1beta1.ReplicaSet](#)
- [v1beta1.ReplicaSetList](#)

Top Level API Objects

- [v1.Service](#)
- [v1.ServiceList](#)
- [v1.Event](#)
- [v1.EventList](#)
- [v1.Namespace](#)
- [v1.NamespaceList](#)
- [v1.Secret](#)
- [v1.SecretList](#)
- [v1.ConfigMap](#)
- [v1.ConfigMapList](#)

Federation v2 -- Concept

The following terms are defined in the context of Federation:

- **Member Cluster:** Federated APIs manage resources in *member clusters*, which are Kubernetes clusters that the Federation API has been configured to access.
- **Template:** A *template* specifies a base Kubernetes resource that should be created in one or more member clusters.
- **Overrides:** *Overrides* enable fields of a template resource to vary by cluster.
- **Placement:** *Placement directives* identify which member clusters a resource or set of resources should exist in.
- **Propagation:** *Propagation* ensures that a resource definition (composed from the template, overrides and placement) exists in a given member cluster.
- **Status:** *Status* provides visibility into the underlying member cluster state that has resulted from propagation (e.g. how many pods exist for a replica set).
- **Scheduling:** *Scheduling* determines resource placement and overrides from both user intent (including per-resource scheduling configuration) and the live state (both spec and status) of member clusters.

Federation v2 -- Base Federated Resources

The pattern for constructing a federated API from an existing k8s API resource:

- **New API group** = federation. + existing API group
- **ResourceNames** = Federated + existing resource name
- **FederatedX's** spec looks like:
 - **template** : the 'template' for what the resource looks like,
absent any cluster-specific differentiation
 - **overrides** : a list of partial overrides of the template for a specific cluster

Federation v2 -- Auxiliary Resources

1. **Cluster Placement Preferences** – a set of hard or soft constraints influencing which clusters the resource should be federated to.
3. **Template Substitution Preferences** – a set of rules defining how and what values are substituted into templates on a per-cluster basis.
4. **Template Substitution Outputs** – These are actual kubernetes API resources, the results of applying cluster placement decisions and template substitution preferences to base federated resource templates.
5. **Resource Propagation Configurations** – these configure the mechanism by which template substitution decisions are propagated to specific clusters,
6. **Propagation Status** ... TBD.

Federation v2 -- Work in progress

```
apiVersion: federation.apps.k8s.io/v1alpha1
kind: FederatedReplicaSet
```

```
metadata:
  name: frontend
  namespace: my-ns
```

```
spec:
  template:
```

```
    apiVersion: apps/v1
    kind: ReplicaSet
    meta:
```

```
      name: frontend
      namespace: my-ns
      metadata:
```

```
        name: nginx
        labels:
          app: frontend
```

```
    spec:
      replicas: 3
      template:
        metadata:
          labels:
            app: guestbook
            tier: frontend
```

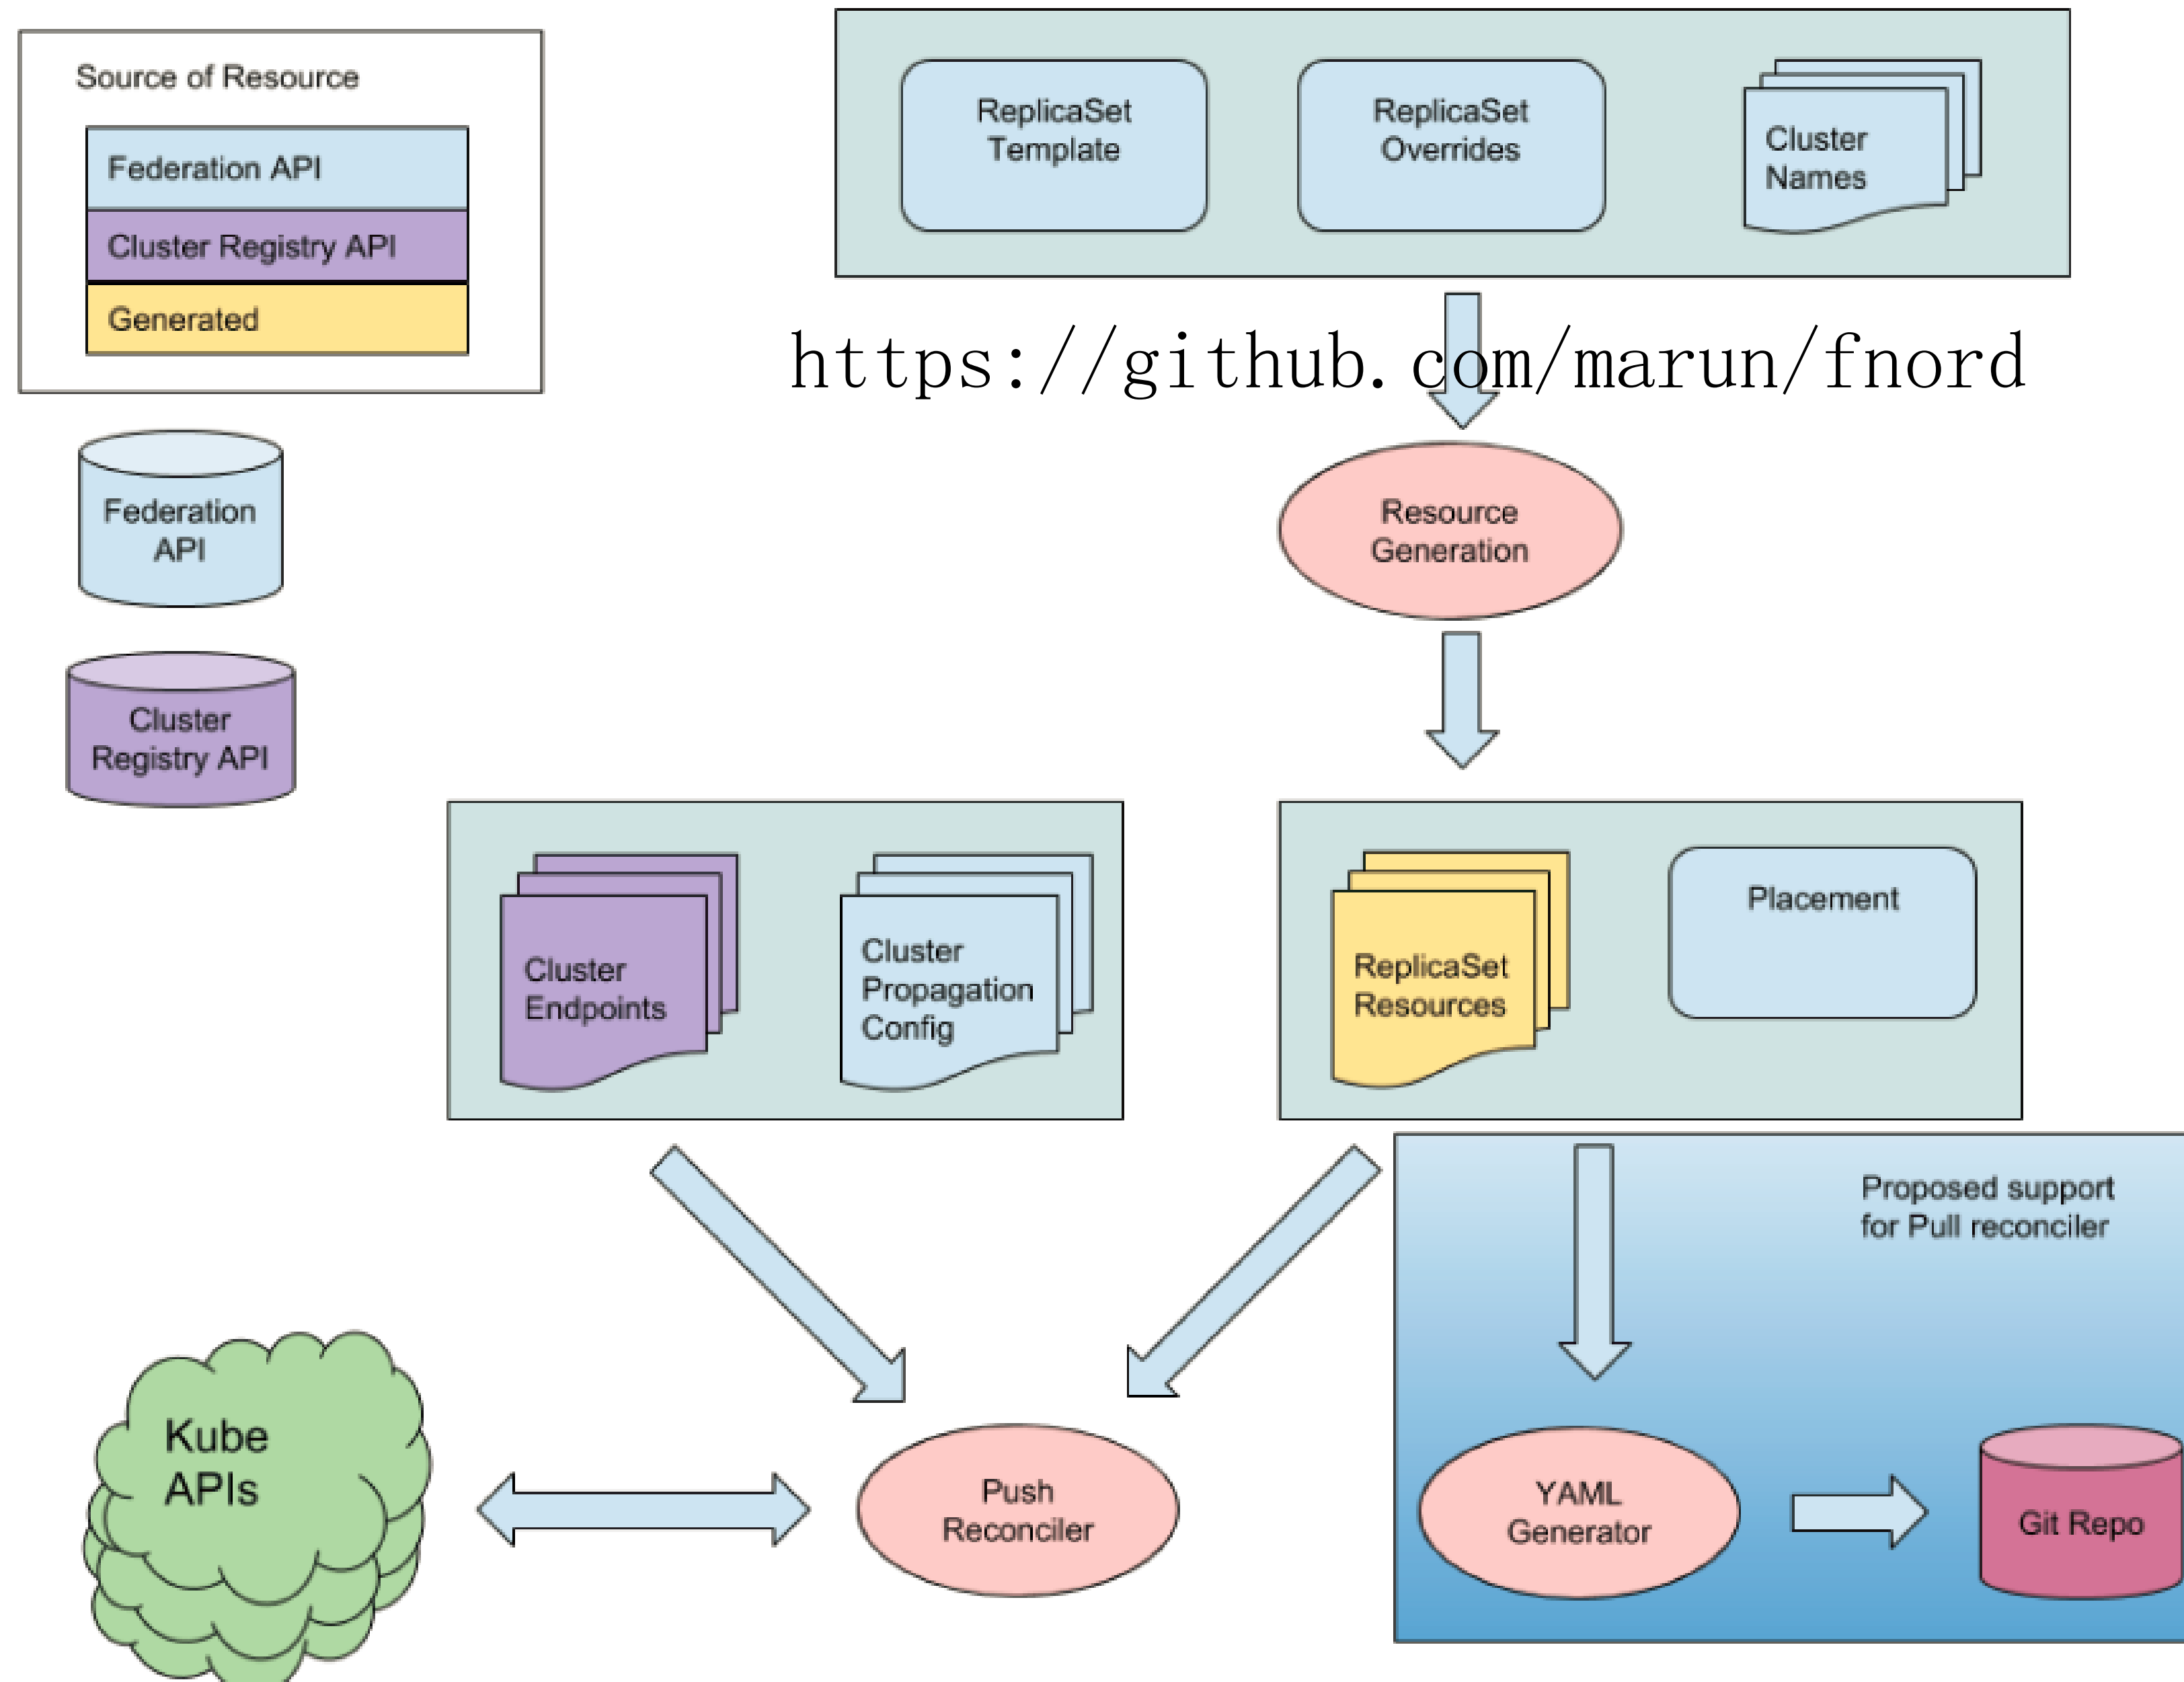
```
        spec:
          containers:
            - name: php-redis
              image: gcr.io/google_samples/gb-frontend:v3
```

```
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v3
          resources:
            requests:
              cpu: 100m
              memory: 100Mi
          env:
            - name: GET_HOSTS_FROM
              value: dns
          ports:
            - containerPort: 80
```

```
    overrides:
      - clusterName: X
        override:
          replicas: 10
          spec:
```

```
            containers:
              - name: php-redis
                image: gcr.io/google_samples/gb-frontend:v4
```

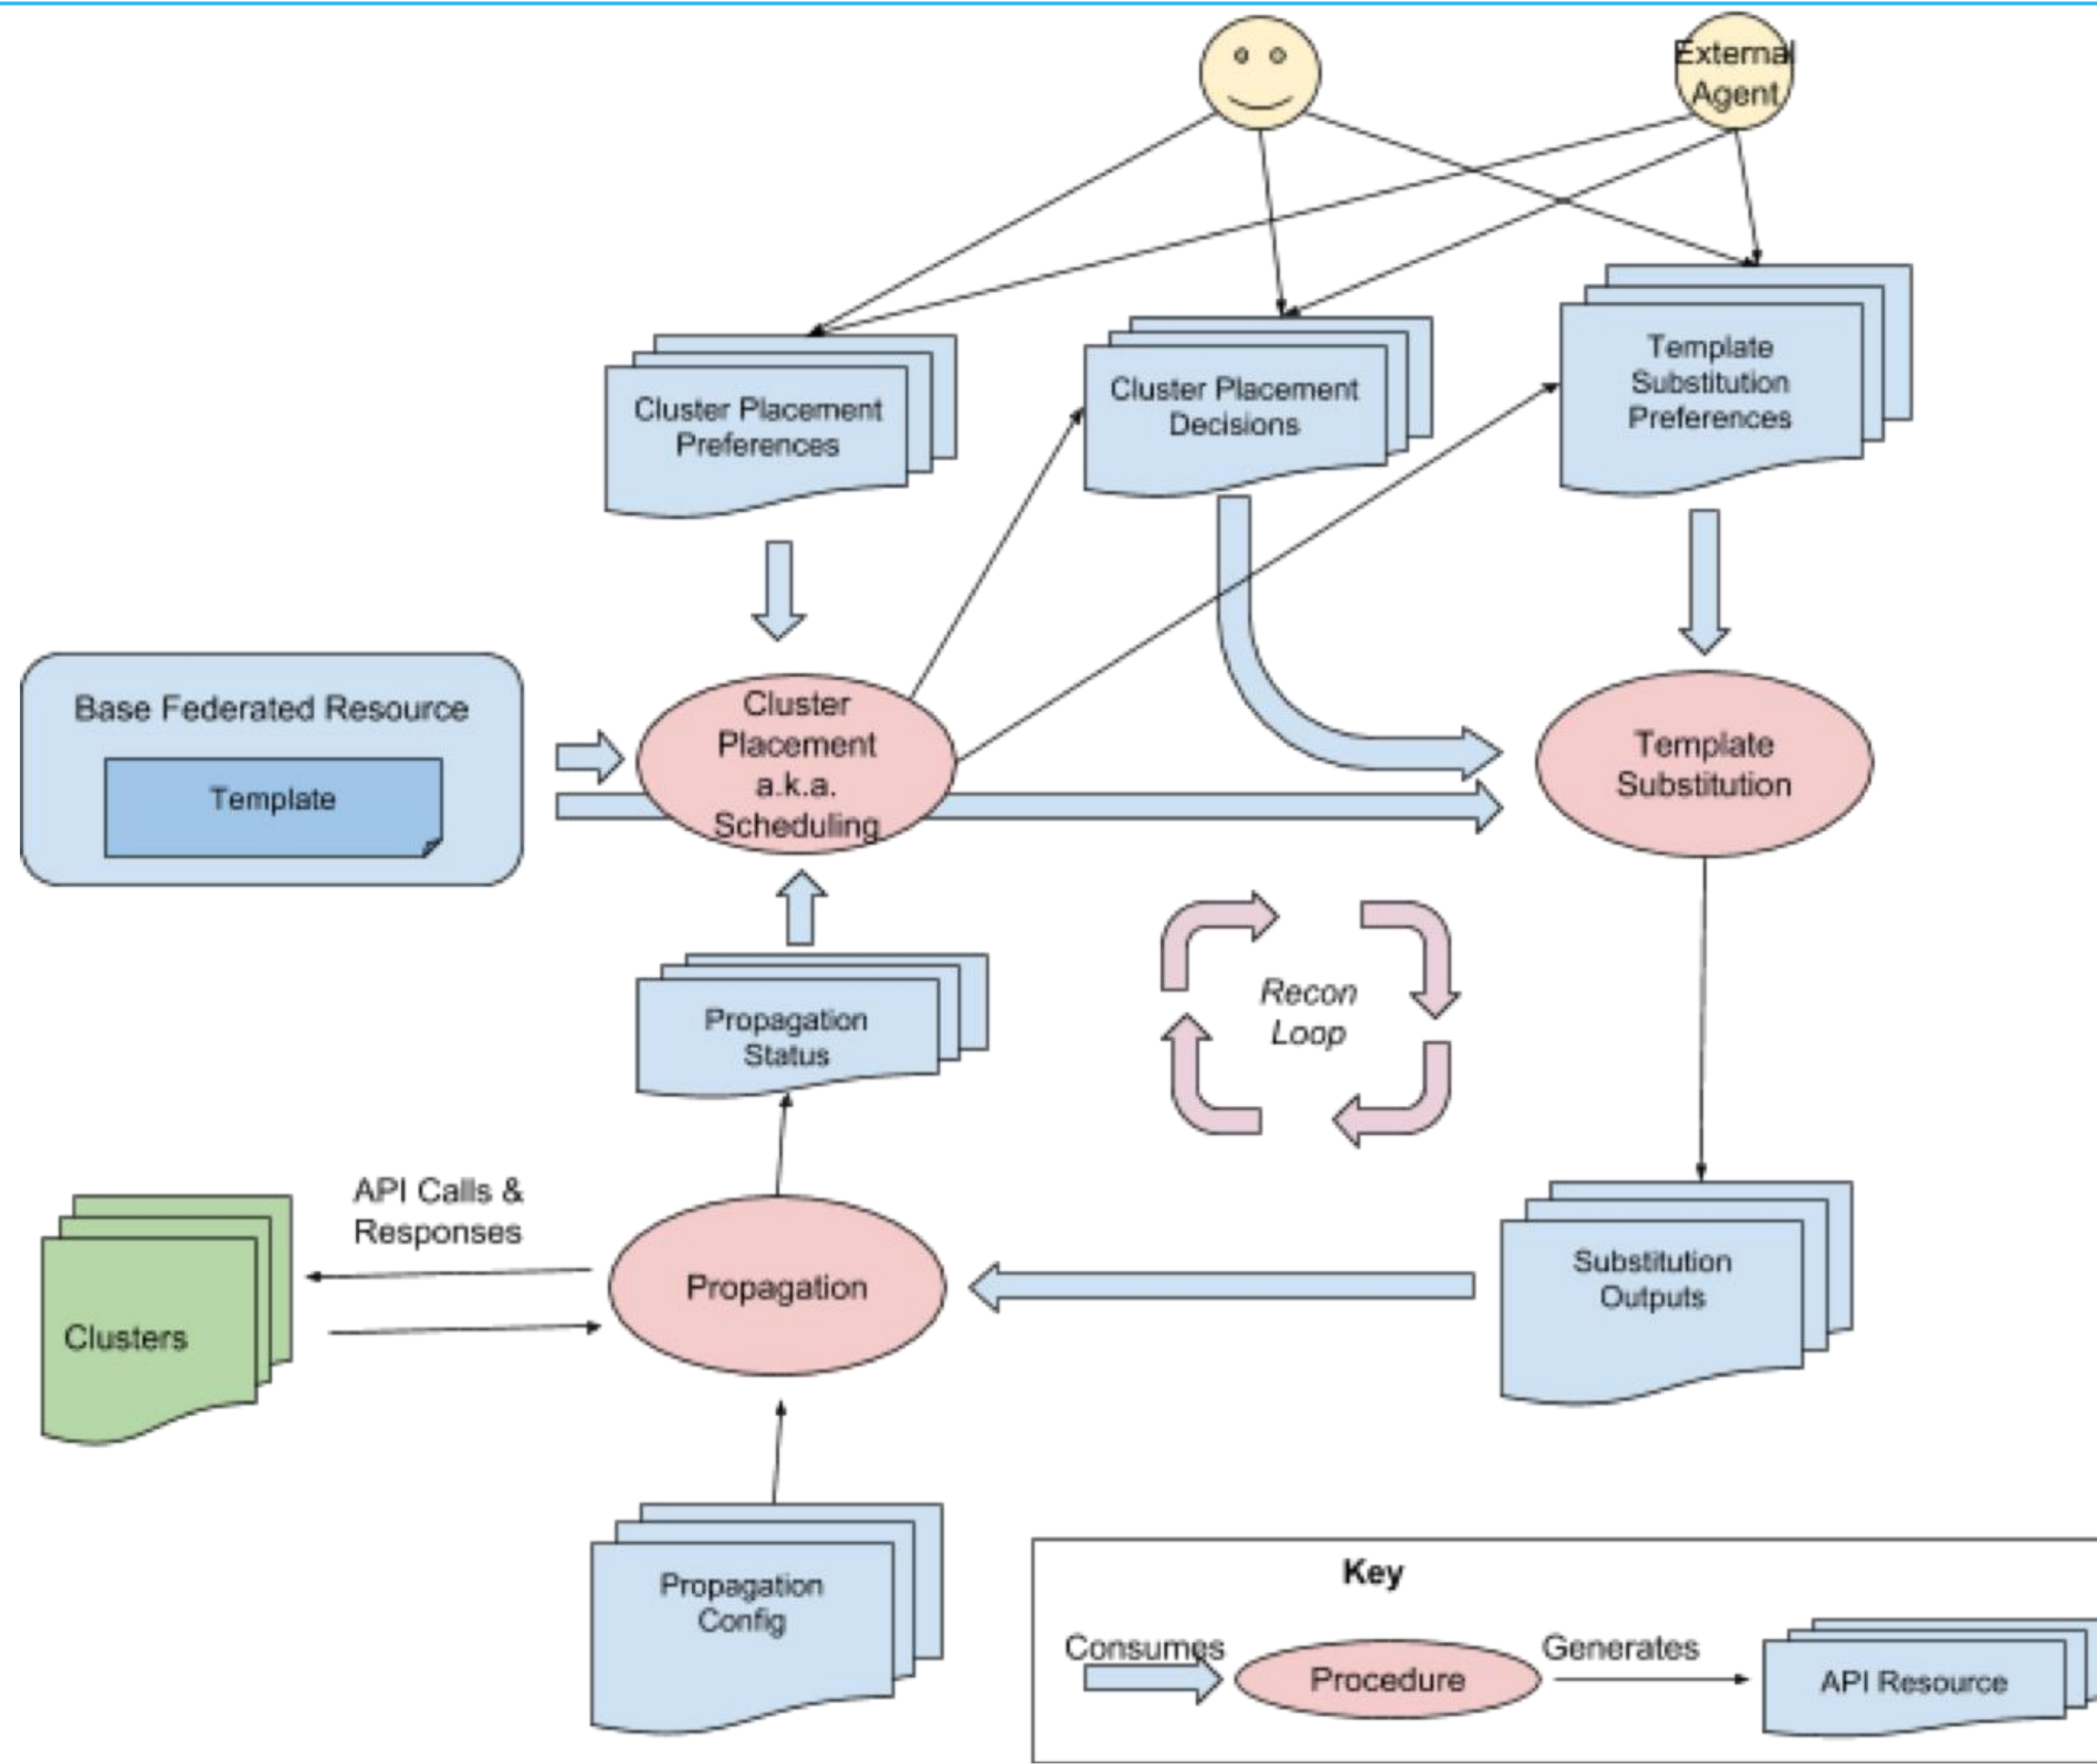

Federation v2 -- Prototype



<https://github.com/marun/fnord>

Maru Newby @Red Hat

Federation v2 -- Work in progress



Global Load balancing

Administrators



End Users





federation-apiserver



federation-controller

app

app



Load balancing



Global Load balancing

Administrators






federation-apiserver


federation-controller

app



app



End Users



Load balancing



Global Load balancing -- Clarifying questions

How does global external service discovery work?

How does dynamic load balancing across clusters ?

How does this tie in with auto-scaling of services?

Global Load balancing

Options for global loadbalancing

Google Global Load Balancer

Ideal solution with multiple clusters on Google Cloud

No solution for Hybrid Cloud (yet)

Backplane.io

Developer focused, very new...

Working Kubernetes and Docker Swarm integration

DNS based (e. g. AWS Route53)

Usual “DNS Problems”, e. g. misbehaving clients ignoring TTLs

Failover times dependent on DNS TTLs and intermediate caching

Content Delivery Networks

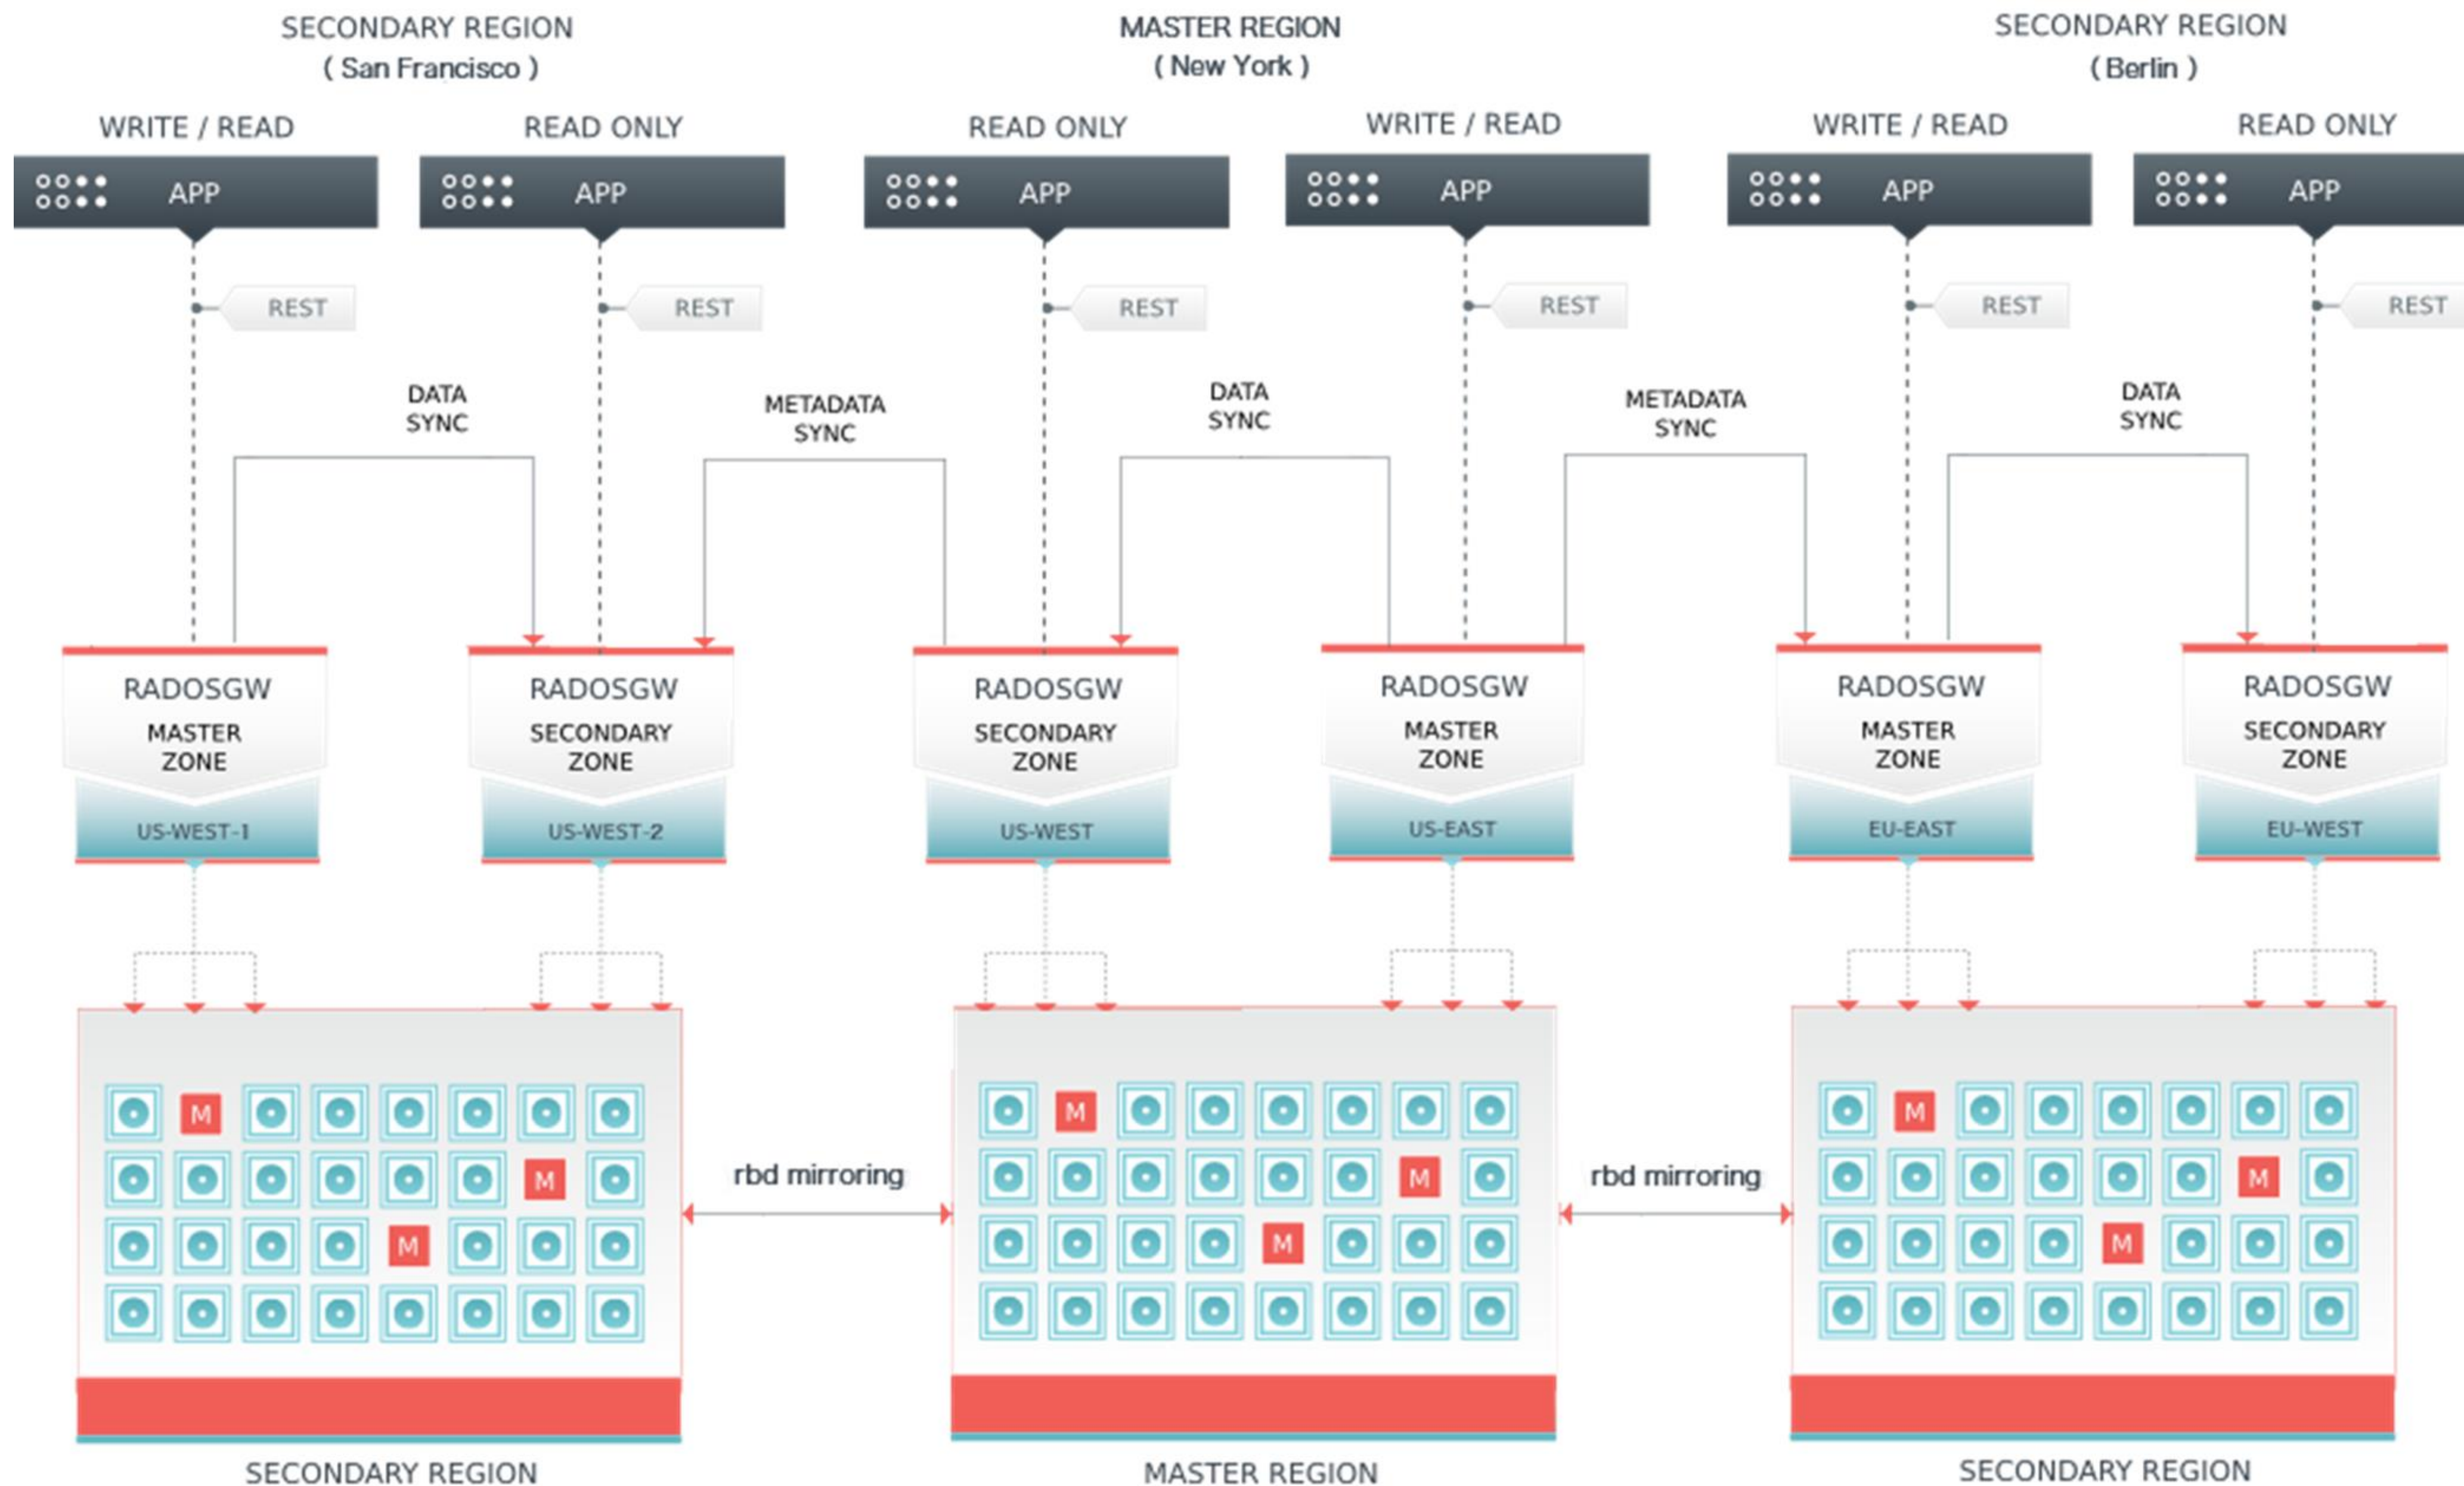
Currently, manual configuration needed

Best option for Hybrid installations now

Global Distributed Storage

1. Ceph
2. Glusterfs
3. Lustre
4. Moosefs
5. HDFS
6.

Global Distributed Storage



Global Distributed Storage

globally distributed databases

like CockroachDB 、 Google Spanner etc.

non-globally distributed databases

like zookeeper、etcd、memcache、redis etc.

Federation -- Caveats & Limitations

1. Still under heavy development
2. Not all functionality is currently available in Federation
3. Single federated control plane
4. Increased network bandwidth and cost
5. A bug in the federation control plane could impact all clusters

Federation – Github Repo



SIG-Federation

Federation v1 (Kubernetes v1.9+)

<https://github.com/kubernetes/federation>

Federation v2

<https://github.com/kubernetes-incubator/apiserver-builder>

<https://github.com/kubernetes/cluster-registry>

<https://github.com/kubernetes-sigs/federation-v2>

<https://github.com/GoogleCloudPlatform/k8s-multicloud-ingress>

<https://github.com/marun/fnord>

Questions ?

